

Practical Machine Learning Class Project

Nick Orka

3/3/2020

Overview

One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. A goal of the project is to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants and to predict the manner in which they did the exercise. This is the “classe” variable in the training set.

Data preprocessing

The training data for this project are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>)

The test data are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>
(<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)

The data for this project come from this source:

<http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>
(<http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har>).

```
## Registered S3 method overwritten by 'GGally':  
##   method from  
##   +.gg      ggplot2
```

Data summary

The initial data have NA data that required to be cleaned:

```
dim(training)
```

```
## [1] 19622  160
```

```
dim(testing)
```

```
## [1]  20 160
```

Cleaning data

Cleaned data:

```

var_names <- names(training)
naNames <- names(training[, colSums(is.na(training)) > 0])
var_names <- var_names[!var_names %in% naNames]
var_names <- var_names[!grepl("X|timestamp|window", var_names)]
training <- training[, var_names]
classe <- training$classe
training <- training[, sapply(training, is.numeric)]
var_names <- names(training)
training$classe <- classe
testing <- testing[, var_names]
dim(training)

```

```
## [1] 19622    53
```

```
dim(testing)
```

```
## [1] 20 52
```

Create validation dataset

```

set.seed(34562)
isTrain <- createDataPartition(training$classe, p = .7, list = F)
validating <- training[!isTrain, ]
training <- training[isTrain, ]
dim(validating)

```

```
## [1] 5885    53
```

```
dim(training)
```

```
## [1] 13737    53
```

Prediction model

```

control <- caret::trainControl(method = "cv", 5)
model <- caret::train(classe ~ ., data = training, method = "rf", trControl = control, ntree = 2
50, localImp = TRUE)
model

```

```
## Random Forest
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10989, 10990, 10989, 10991, 10989
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa
##    2    0.9900269  0.9873832
##   27    0.9909734  0.9885804
##   52    0.9842034  0.9800127
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

Then, we estimate the performance of the model on the validation data set.

```
p <- predict(model, validating)
confusionMatrix(validating$classe, p)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1673    1    0    0    0
##           B   11 1128    0    0    0
##           C    0    9 1015    2    0
##           D    0    0   15  948    1
##           E    0    1    3    6 1072
##
## Overall Statistics
##
##           Accuracy : 0.9917
##           95% CI : (0.989, 0.9938)
##           No Information Rate : 0.2862
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9895
##
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9935  0.9903  0.9826  0.9916  0.9991
## Specificity      0.9998  0.9977  0.9977  0.9968  0.9979
## Pos Pred Value   0.9994  0.9903  0.9893  0.9834  0.9908
## Neg Pred Value   0.9974  0.9977  0.9963  0.9984  0.9998
## Prevalence       0.2862  0.1935  0.1755  0.1624  0.1823
## Detection Rate   0.2843  0.1917  0.1725  0.1611  0.1822
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9966  0.9940  0.9902  0.9942  0.9985
```

```
accuracy <- postResample(p, validating$classe)
accuracy
```

```
## Accuracy      Kappa
## 0.9916737 0.9894659
```

```
oose <- 1 - as.numeric(confusionMatrix(validating$classe, p)$overall[1])
oose
```

```
## [1] 0.008326253
```

So, the estimated accuracy of the model is 99.17% and the estimated out-of-sample error is 0.83%.

Predicting for Test Data Set

Now, we apply the model to the original testing data set downloaded from the data source.


```
result <- predict(model, testing)
result
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Appendix: Figures


1. Variable importance

Variable Importance of RF model

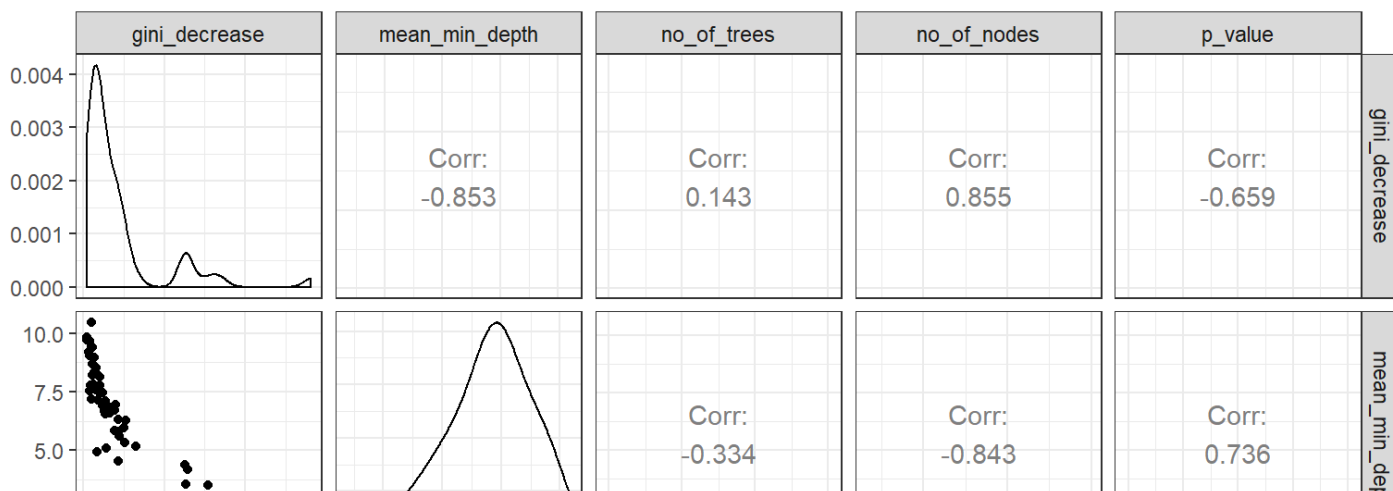


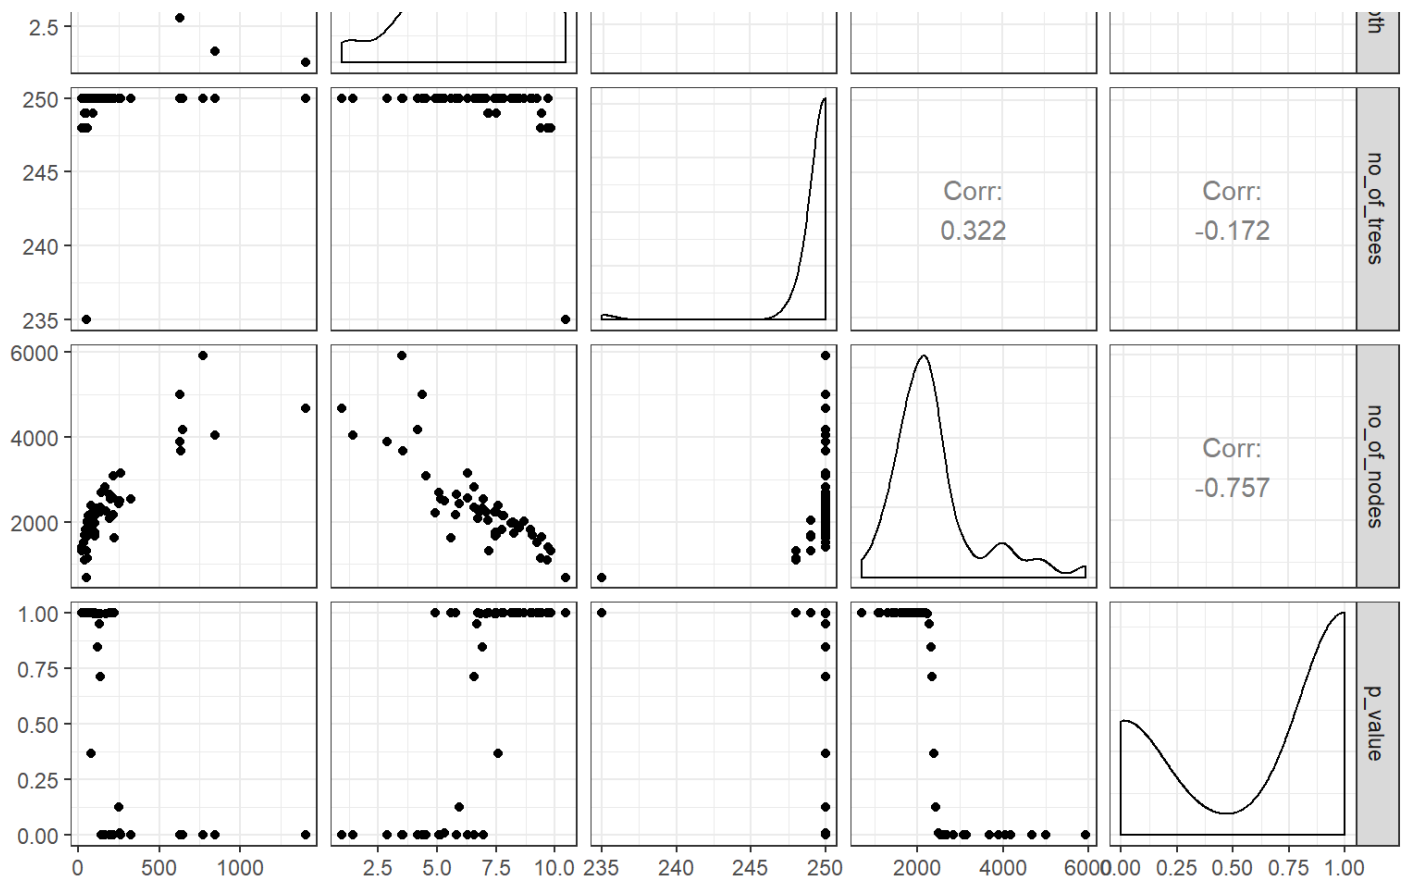
```
## Warning: Using alpha for a discrete variable is not advised.
```

Multi-way importance plot




Relations between measures of importance





2. Correlation Matrix Visualization



3. Decision Tree Visualization

